

ISIT312 Big Data Management

Hadoop Architecture

Dr Fenghui Ren

School of Computing and Information Technology -
University of Wollongong

Hadoop Architecture

Outline

Hadoop Distributed File System (HDFS)

NameNode metadata

DataNode and Secondary node

Yet Another Resource Negotiator (YARN)

ResourceManger

NodeManager

ApplicationMaster

Summary

HDFS: Hadoop Distributed File System

HDFS is designed for:

- Very large files
- Stream data access
- Commodity hardware

But not for:

- Low-latency data access
- Lots of small files
- Multiple writers, arbitrary file modifications

HDFS: Hadoop Distributed File System

HDFS contains the following key components:

NameNode:

- HDFS master node process
- manages the filesystem metadata
- does not store a file itself

SecondaryNameNode and Standby NameNode

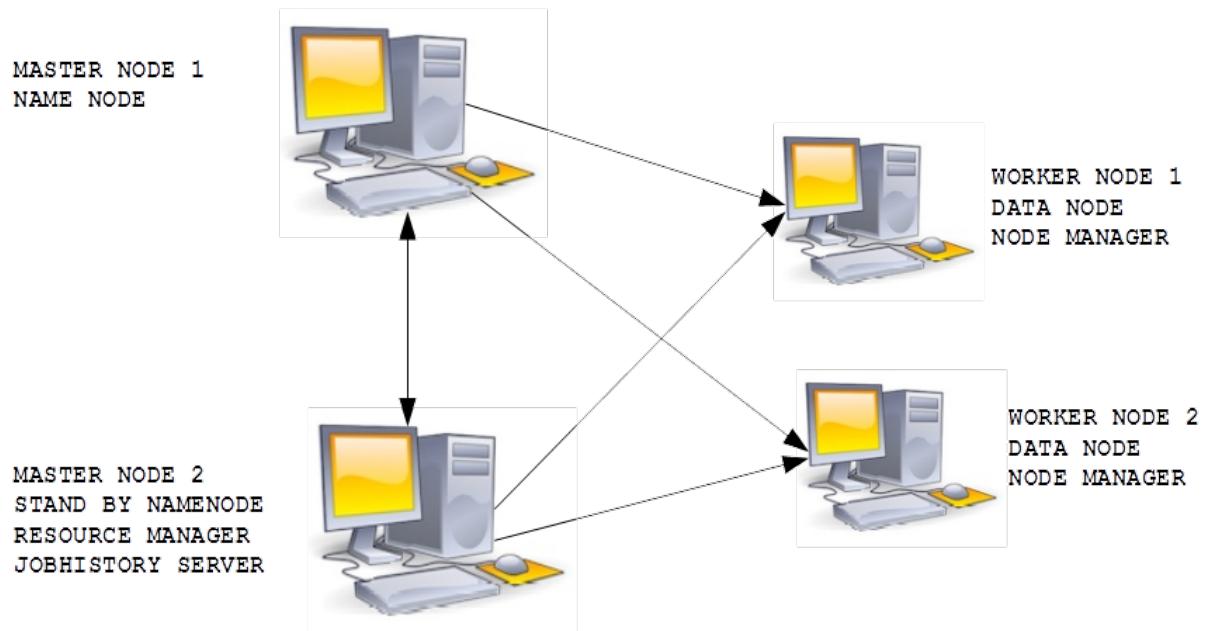
- SecondaryNameNode expedites the filesystem metadata recovery
- Standby NameNode (optional) provides high availability

DataNode

- runs HDFS slave node process
- manages block storage and access for reading or writing of data, block replication

HDFS: Hadoop Distributed File System

Architecture of HDFS



HDFS: Hadoop Distributed File System

HDFS is a virtual filesystem

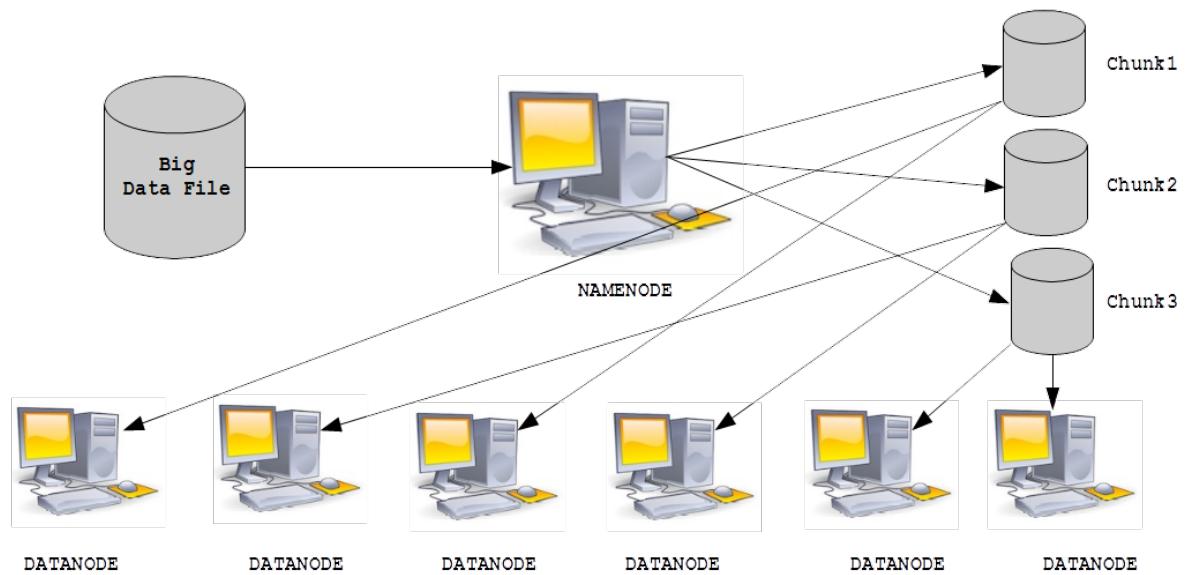
- appears to a client as one file system, but the data is stored in multiple different locations
- deployed on the top of the native filesystems (such as `ext3`, `ext4` and `xfs` in Linux)

Each file in HDFS consists of blocks

- The size of each block defaults to 128MB but is configurable
- The default number of replicates for blocks is 3, but it is also configurable

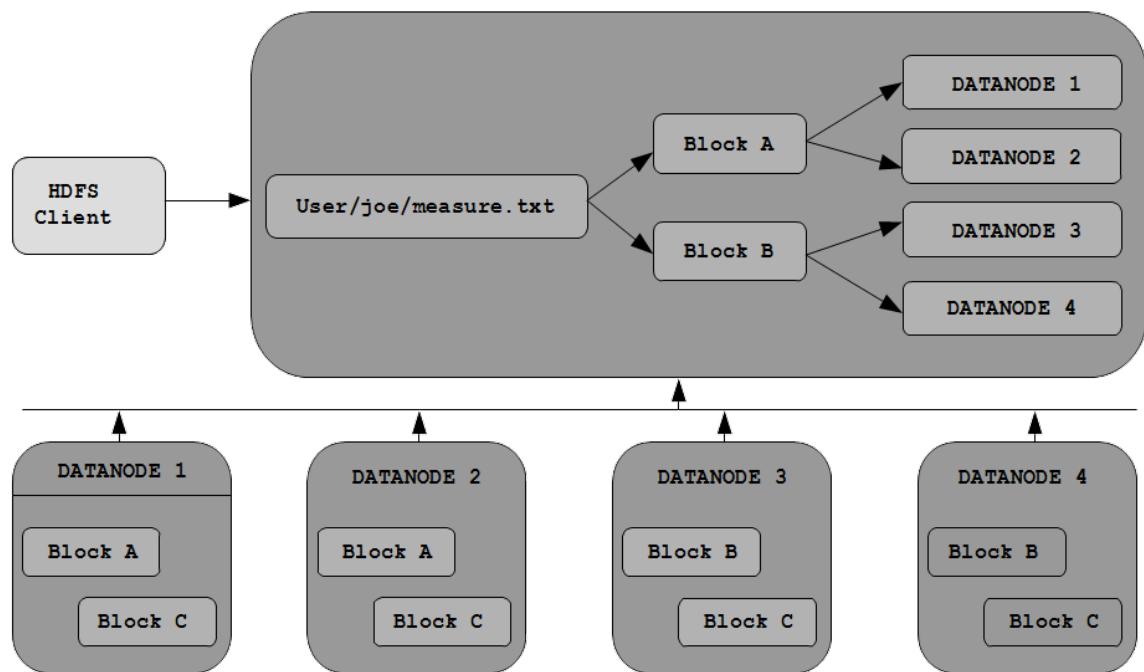
HDFS: Hadoop Distributed File System

Logical view of data storage



HDFS: Hadoop Distributed File System

Physical implementation of data file storage



Hadoop Architecture

Outline

[Hadoop Distributed File System \(HDFS\)](#)

[NameNode metadata](#)

[DataNode and Secondary node](#)

[Yet Another Resource Negotiator \(YARN\)](#)

[ResourceManger](#)

[NodeManager](#)

[ApplicationMaster](#)

[Summary](#)

NameNode Metadata

NameNode stores the metadata of the files in HDFS

| object | block_id | seq | locations | ACL | Checksum |
|----------------|-----------|-----|---------------------|------------|--------------|
| /data/file.txt | blk_00123 | 1 | [node1,node2,node3] | -rwxrwxrwx | 8743b52063.. |
| /data/file.txt | blk_00124 | 2 | [node2,node3,node4] | -rwxrwxrwx | cd84097a65.. |
| /data/file.txt | blk_00125 | 3 | [node2,node4,node5] | -rwxrwxrwx | d1633f5c74.. |

NameNode functions:

- Maintain the metadata pertaining to the file system (e.g., the file hierarchy and the block locations for each file)
- Manage user access to the data files
- Map the data blocks to the DataNodes in the cluster
- Perform file system operations (e.g., opening and closing the files and directories)
- Provide registration services and periodic heartbeats for DataNodes

Hadoop Architecture

Outline

[Hadoop Distributed File System \(HDFS\)](#)

[NameNode metadata](#)

[DataNode and Secondary node](#)

[Yet Another Resource Negotiator \(YARN\)](#)

[ResourceManger](#)

[NodeManager](#)

[ApplicationMaster](#)

[Summary](#)

DataNode and Secondary node

DataNode functions:

- Provide the block storage by storing blocks on the local file system
- Fulfil the read/write requests
- Replicating data across the cluster
- Keeping in touch with the [NameNode](#) by sending periodic block reports and heartbeats
- A heartbeat confirms the [DataNode](#) is alive and healthy, and a block report shows the blocks being managed by the [DataNode](#)

Secondary NameNode and Standby NameNode functions:

- Without a [NameNode](#), there is no way to know to which files the blocks stored on the [DataNodes](#) correspond to
- In essence, all files in [HDFS](#) are lost
- [Secondary NameNode](#) periodically backups the metadata in the (primary) [NameNode](#), which is usually for recovery
- [Standby NameNode](#) is a hot node that runs together with the (primary) [NameNode](#) in the cluster, facilitating high-availability

Hadoop Architecture

Outline

[Hadoop Distributed File System \(HDFS\)](#)

[NameNode metadata](#)

[DataNode and Secondary node](#)

[Yet Another Resource Negotiator \(YARN\)](#)

[ResourceManger](#)

[NodeManager](#)

[ApplicationMaster](#)

[Summary](#)

Yet Another Resource Negotiator (YARN)

[YARN](#): the core subsystem in Hadoop responsible for governing, allocating, and managing the finite distributed processing resources available on a Hadoop cluster

- introduced in Hadoop 2 to improve the [MapReduce](#) implementation, but general enough to support other distributed computing paradigms

[YARN](#) provides its core services via two types of long-running daemons:

- A [ResourceManager](#) (one per cluster) to manage the use of resources across the cluster, and
- [NodeManagers](#) running on all the nodes in the cluster to launch and monitor containers

Yet Another Resource Negotiator (YARN)

Architecture of YARN

A **client** is the program that submits jobs to the cluster

- May also be the gateway machine that the client program runs on

A **job**, also called an **application**, contains one or more tasks

- A task in a MapReduce job can be either a **mapper** and a **reducer task**

Each **mapper** and **reducer** task runs within a **container**

- **Containers** are logical constructs that represent a specific amount of memory and other resources, such as processing cores (CPU)
- For example, a **container** can represent 2GB memory and 2 processing cores
- **Containers** may also refer to the running environment of an application

Yet Another Resource Negotiator (YARN)

Architecture of YARN

[ResourceManager](#): YARN's daemon running on a master node

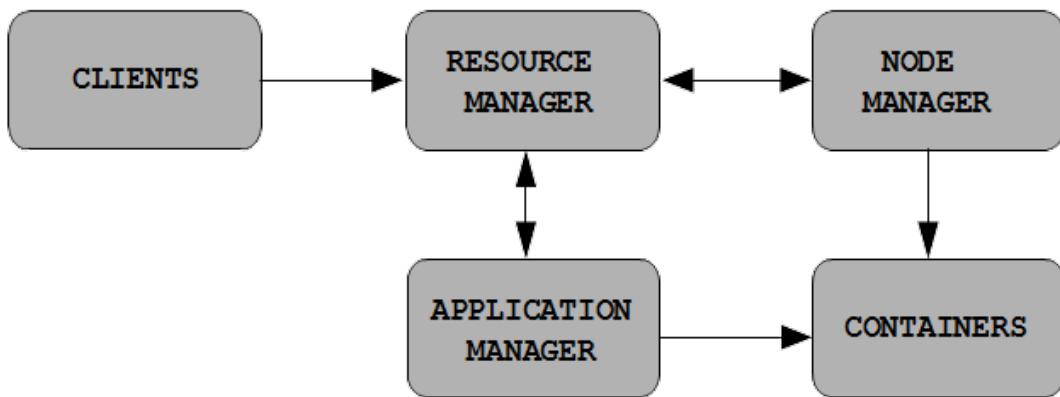
- [ResourceManager](#) is responsible for granting cluster computing resources to applications running on the cluster
- Resources are granted the items of containers

[NodeManager](#): YARN's daemon running on a slave node.

- [NodeManager](#) manages containers on a slave node
- [ApplicationMaster](#): the first container allocated by the [ResourceManager](#) to run on a [NodeManager](#) for each application

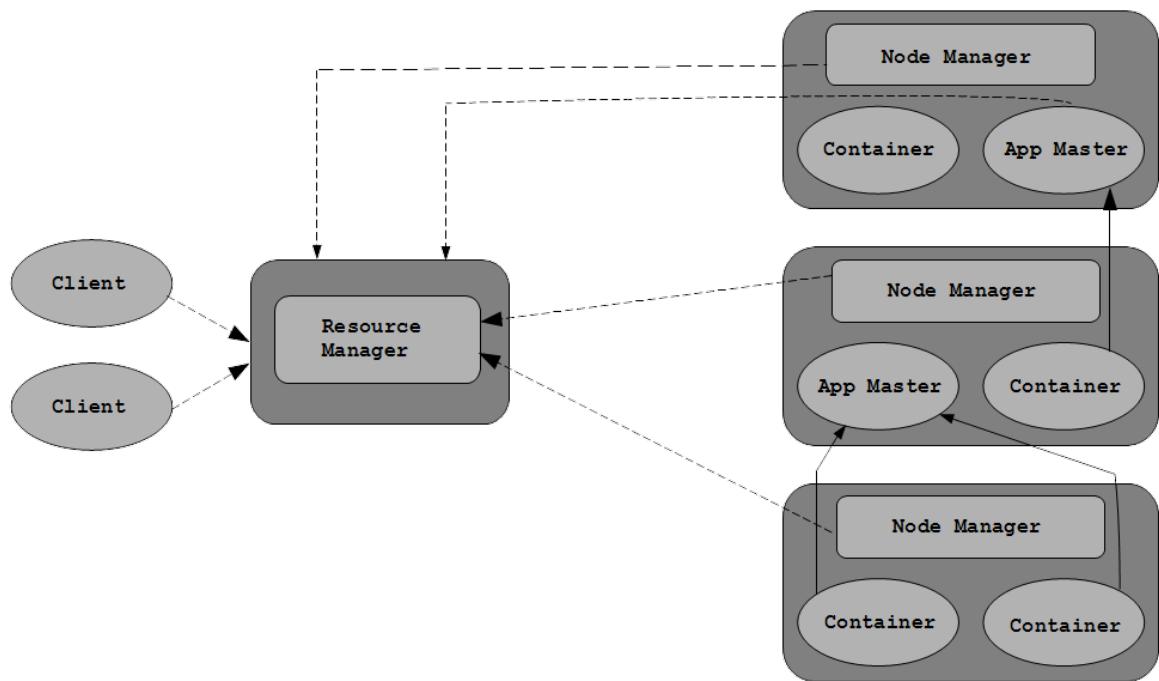
Yet Another Resource Negotiator (YARN)

Architecture of YARN



Yet Another Resource Negotiator (YARN)

Architecture of YARN



Hadoop Architecture

Outline

[Hadoop Distributed File System \(HDFS\)](#)

[NameNode metadata](#)

[DataNode and Secondary node](#)

[Yet Another Resource Negotiator \(YARN\)](#)

[ResourceManger](#)

[NodeManager](#)

[ApplicationMaster](#)

[Summary](#)

ResourceManager

There is one [ResourceManager](#) per cluster, which consists of two key components: [Scheduler](#) and [ApplicationManager](#)

Key functions of [ResourceManager](#):

- Creates the first container for an application to run [ApplicationMaster](#) for that application
- Tracks the heartbeats from [NodeManagers](#) to manage [DataNodes](#)
- Runs [Scheduler](#) to determine resource allocation among the clusters
- Manages cluster level security
- Manages the resource requests from [ApplicationMasters](#)
- Monitors the status of [ApplicationMasters](#) and restarts that container upon its failure
- Deallocates the containers when the application completes or after they expire

The role of [ResourceManager](#) is pure management and scheduler

It does not perform any actual data processing, for example the **Map** and **Reduce** functions in a [MapReduce](#) application

Hadoop Architecture

Outline

[Hadoop Distributed File System \(HDFS\)](#)

[NameNode metadata](#)

[DataNode and Secondary node](#)

[Yet Another Resource Negotiator \(YARN\)](#)

[ResourceManger](#)

[NodeManager](#)

[ApplicationMaster](#)

[Summary](#)

NodeManager

Each [DataNode](#) runs a [NodeManager](#) daemon for performing [YARN](#) functions

Main functions of a [NodeManager](#) daemon:

- Communicates with [ResourceManager](#) through health heartbeats and container status notifications.
- Registers and starts the application processes
- Launches both [ApplicationMaster](#) and the rest of an application's resource containers (that is, the map and reduce tasks that run in the containers) on request from [ApplicationMaster](#)
- Oversees the lifecycle of the application containers
- Monitors, manages and provides information regarding the resource consumption (CPU/memory) by the containers
- Tracks the health of [DataNode](#)
- Provides auxiliary services to [YARN](#) applications, such as services used by the MapReduce framework for its shuffle and sort operations

Hadoop Architecture

Outline

[Hadoop Distributed File System \(HDFS\)](#)

[NameNode metadata](#)

[DataNode and Secondary node](#)

[Yet Another Resource Negotiator \(YARN\)](#)

[ResourceManger](#)

[NodeManager](#)

[ApplicationMaster](#)

[Summary](#)

ApplicationMaster

For each **YARN** application, there is a dedicated **ApplicationMaster**

Functions of **ApplicationMaster**:

- Managing task scheduling and execution
- Allocating resources locally for the application's tasks

ApplicationMaster is running within a container

ApplicationMaster's existence is associated with the running application

When an application is completed, its **ApplicationMaster** no longer exists

Once created, **ApplicationMaster** is in charge of requesting resources with **ResourceManager** to run the application

The resource request are very specific, for example:

- the file blocks needed to process the job,
- the amount of the resource, in terms of the number of containers to create for the application,
- the size of the containers, etc.

Hadoop Architecture

Outline

[Hadoop Distributed File System \(HDFS\)](#)

[NameNode metadata](#)

[DataNode and Secondary node](#)

[Yet Another Resource Negotiator \(YARN\)](#)

[ResourceManger](#)

[NodeManager](#)

[ApplicationMaster](#)

[Summary](#)

Summary

Terminologies

*For convenience, we use the **names of HDFS and YARN processes** to refer to both the **hosts** and the **daemons** running on the corresponding hosts*

*For example, **RecourseManager** refers to both a **master node** and the **RecourseManager daemon** on that master node; **DataNode** refers to both a **slave node** and the **DataNode daemon** on that slave node.*

Hadoop is a leading platform for big data

Hadoop consists of a storage layer (**HDFS**), a coordination and management layer (**YARN**) and a processing layer (e.g., **MapReduce**)

HDFS and **YARN** have key services (daemons)

MapReduce is a fundamental computing model (i.e., batch processing) for big data

Next: Interaction with Hadoop and "dive" into the **MapReduce** framework

References

White T., Hadoop The Definitive Guide: Storage and analysis at Internet scale, O'Reilly, 2015 (Available through UOW library)

Vohra D., Practical Hadoop ecosystem: a definitive guide to Hadoop-related frameworks and tools, Apress, 2016 (Available through UOW library)

Aven J., Hadoop in 24 Hours, SAMS Teach Yourself, SAMS 2017

Alapati S. R., Expert Hadoop Administration: Managing, tuning, and securing Spark, YARN and HDFS, Addison-Wesley 2017