

ISIT312 Big Data Management

Spark Operations

Dr Fenghui Ren

School of Computing and Information Technology -
University of Wollongong

Spark Operations

Outline

The Programming Language Scala

Quick Start

Self Contained Application

Web User Interface

Operations on Resilient Distributed Datasets (RDDs)

Operations on Datasets

Operations on DataFrames

SQL Module

The Programming Language Scala

Spark has built-in APIs for Java, Scala, and Python, and is also integrated with R

Among all languages, Scala is the most supported language

Also, Spark project is implemented using Scala

Therefore, we choose Scala as our working language in Spark

Scala is a Java-like programming language which unifies object-oriented and functional programming

Scala is a pure object-oriented language in the sense that every value is an object

Types and behaviour of objects are described by classes

Scala is a functional programming language in the sense that every function is a value

Nesting of function definitions and higher-order functions are naturally supported

The Programming Language Scala

Hello World ! in Scala

```
object Hello {  
    def main(args: Array[String]) = {  
        println("Hello, world")  
    }  
}
```

Hello World ! in Scala

Instead of including `main` method, it can be extended with `App` trait

```
object Hello2 extends App {  
    println("Hello, world")  
}
```

Extending App trait

Using command line arguments

```
object HelloYou extends App {  
    if (args.size == 0)  
        println("Hello, you")  
    else  
        println("Hello, " + args(0))  
}
```

Command line arguments

The Programming Language Scala

Difference between `var`, `val`, and `def`

```
var x = 7  
x = x * 2
```

Variable

```
val x = 7  
x = x * 2  
'error: reassignment to val'
```

Value

```
def hello(name: String) = "Hello : " + name  
hello("James") // "Hello : James"  
hello("") // "Hello : "
```

Function declaration

When `lazy` keyword is used then a value is only computed when it is needed

```
lazy val x = {  
    println("calculating value of x")  
    13 }  
val y = {  
    println("calculating value of y")  
    20 }
```

Lazy evaluation

The Programming Language Scala

Defining a class

```
class Point(var x: Int, var y: Int) {  
  
    def move(dx: Int, dy: Int): Unit = {  
        x = x + dx  
        y = y + dy  
    }  
  
    override def toString: String =  
        s"($x, $y")  
}  
  
val point1 = new Point(2, 3)  
println(point1.x)           // 2  
println(point1)            // prints (2, 3)
```

Class Point

Method move

Method toString

Applications

Spark Operations

Outline

The Programming Language Scala

Quick Start

Self Contained Application

Web User Interface

Operations on Resilient Distributed Datasets (RDDs)

Operations on Datasets

Operations on DataFrames

SQL Module

Quick Start

To open [Scala](#) version of [Spark shell](#) in standalone mode process the following command

```
bin/spark-shell --master local[*]
```

Starting Spark shell in standalone mode

To open **Spark shell** shell with **YARN**, process the following command

```
bin/spark-shell --master yarn
```

Starting Spark shell with Yarn

Quick Start

A **SparkSession** instance is an entry to a **Spark** application

- If you type **spark** in the spark-shell interface then you get the following messages

```
res0: org.apache.spark.sql.SparkSession = org.apache.spark.sql.SparkSession@...
```

Message

You can use **SparkSession** instance **spark** to interact with **Spark** and to develop your data processing pipeline

For example,

```
val myRange = spark.range(1000).toDF("number")
myRange: org.apache.spark.sql.DataFrame = [number: bigint]
```

Creating Data Frame

```
myRange.show(2)
+----+
| number |
+----+
|     0 |
|     1 |
+----+
```

Listing Data Frame

Quick Start

Sample processing of a file **README.md**

```
Setting Spark Home folder
val YOUR_SPARK_HOME ="path-to-your-Spark-home"

Reading a text file
val textFile = spark.read.textFile("$YOUR_SPARK_HOME/README.md")
textFile: org.apache.spark.sql.Dataset[String] = [value: string]

Counting rows
textFile.count()
res0: Long = 104

Reading the first row
textFile.first()
res1: String = # Apache Spark

Filtering and counting rows
textFile.filter(line => line.contains("Spark")).count()
res2: Long = 20
```

Quick Start

More operations on a file

Counting number of words in the longest line

```
textFile.map(line => line.split(" ").size).reduce((a, b) => if (a > b) a else b)  
res3: Int = 22
```

Filtering and counting rows

```
val wordCounts = textFile.flatMap(line => line.split("")).groupByKey(identity).count()  
wordCounts: org.apache.spark.sql.Dataset[(String, Long)] = [value: string, count(1): bigint]
```

Listing results

```
wordCounts.show(2)  
+-----+-----+  
|    value|count(1)|  
+-----+-----+  
|    online|      1|  
|    graphs|      1|  
+-----+-----+  
only showing top 2 rows
```

Listing results

```
wordCounts.collect()  
res7: Array[(String, Long)] = Array((online,1), (graphs,1), ("Parallel,1), ("Building,1), (thread,1),  
(documentation,3), (command,,2), (abbreviated,1), (overview,1), (rich,1), (set,2), ...)
```

Spark Operations

Outline

The Programming Language Scala

Quick Start

Self Contained Application

Web User Interface

Operations on Resilient Distributed Datasets (RDDs)

Operations on Datasets

Operations on DataFrames

Spark SQL Module

Self-Contained Application

A sample self-contained application

SimpleApp.scala

```
import org.apache.spark.sql.SparkSession
object SimpleApp {
    def main(args: Array[String]) {
        val logFile = "YOUR_SPARK_HOME/README.md"
        // Should be some file on your system
        val spark = SparkSession.builder
            .appName("Simple Application")
            .config("spark.master", "local[*]")
            .getOrCreate()
        val logData = spark.read.textFile(logFile).cache()
        val numAs = logData.filter(line => line.contains("a")).count()
        val numBs = logData.filter(line => line.contains("b")).count()
        println(s"Lines with a: $numAs, Lines with b: $numBs")
        spark.stop()
    }
}
```

Self-Contained Application

Compiling **Scala** source code using **scalac**

```
scalac -classpath "$SPARK_HOME/jars/*" SimpleApp.scala
```

Compiling Scala source code

Creating a jar file in the following way

```
jar cvf app.jar SimpleApp*.class
```

Creating jar

Process it with Spark-shell in the following way

```
$SPARK_HOME/bin/spark-submit --master local[*] --class SimpleApp app.jar
```

Processing

Spark Operations

Outline

The Programming Language Scala

Quick Start

Self Contained Application

Web User Interface

Operations on Resilient Distributed Datasets (RDDs)

Operations on Datasets

Operations on DataFrames

Spark SQL Module

Web UI

Each driver program has a Web UI, typically on port **4040**

Spark Web UI displays information about running tasks, executors, and storage usage.

The screenshot shows the Apache Spark 2.1.0-SNAPSHOT Web UI. The top navigation bar includes links for Jobs, Stages, Storage, Environment, Executors, SQL, and Spark shell application UI. The main content area is titled "Spark Jobs (?)". It displays summary statistics: User: jacek, Total Uptime: 35 s, Scheduling Mode: FIFO, Active Jobs: 1, Completed Jobs: 1, Failed Jobs: 1. Below this is a link to "Event Timeline". The "Active Jobs (1)" section contains a table with one row, showing Job Id 2, Description "show at <console>:24", Submitted 2016/09/29 14:01:20, Duration 5 s, Stages: Succeeded/Total 0/1, and Tasks (for all stages): Succeeded/Total 0/1. The "Completed Jobs (1)" section contains a table with one row, showing Job Id 0, Description "show at <console>:24", Submitted 2016/09/29 14:01:07, Duration 0.3 s, Stages: Succeeded/Total 1/1, and Tasks (for all stages): Succeeded/Total 1/1. The "Failed Jobs (1)" section contains a table with one row, showing Job Id 1, Description "show at <console>:24", Submitted 2016/09/29 14:01:14, Duration 87 ms, Stages: Succeeded/Total 0/1 (1 failed), and Tasks (for all stages): Succeeded/Total 0/1 (1 failed).

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
2	show at <console>:24	2016/09/29 14:01:20	5 s	0/1	0/1

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
0	show at <console>:24	2016/09/29 14:01:07	0.3 s	1/1	1/1

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	show at <console>:24	2016/09/29 14:01:14	87 ms	0/1 (1 failed)	0/1 (1 failed)

Spark Operations

Outline

The Programming Language Scala

Quick Start

Self Contained Application

Web User Interface

Operations on Resilient Distributed Datasets (RDDs)

Operations on Datasets

Operations on DataFrames

Spark SQL Module

Operations on Resilient Distributed Datasets (RDDs)

Operations on **RDDs** are performed on raw **Java** or **Scala** objects

Creating a simple **RDD** with words and distributing over 2 partitions

Creating RDD

```
val myCollection = "Spark The Definitive Guide : Big Data Processing Made Simple".split(" ")
val words = spark.sparkContext.parallelize(myCollection, 2)
```

Eliminating duplicates and counting words

Distinct and counting

```
words.distinct().count()
```

Filtering

Filtering function

```
def startsWithS(individual:String) = { individual.startsWith("S") }
```

Filtering

```
val onlyS = words.filter(word => startsWithS(word))
```

Results of filtering

```
onlyS.collect()
```

TOP

Operations on Resilient Distributed Datasets (RDDs)

Sorting of RDD uses `sortBy` method and a function that extracts a value from the objects

Sorting

```
words.sortBy(word => word.length() * -1).take(2))
```

Random split into Array

Split into Array

```
val fiftyFiftySplit = words.randomSplit(Array[Double](0.5, 0.5))
```

Reduce RDD to one value

Reducing

```
def wordLengthReducer(leftWord:String, rightWord:String): String = {
  if (leftWord.length >= rightWord.length)
    return leftWord
  else
    return rightWord }
```

Reducing

```
words.reduce(wordLengthReducer)
```

Operations on Resilient Distributed Datasets (RDDs)

Some operations on **RDDs** are available on key-value pairs

The most common ones are distributed "shuffle" operations, such as grouping or aggregating the elements by a key

For example, **reduceByKey** operation on key-value pairs can be used to count how many times each line of text occurs in a file

```
val lines = sc.textFile("data.txt")
val pairs = lines.map(s => (s, 1))
val counts = pairs.reduceByKey((a, b) => a + b)
```

Sorting

Some of the transformations of **RDDs**

map(func):	passes each element of RDD through a function	Transformations
filter(func):	selects all element for which a function returns true	
sample(withReplacement, fraction, seed):	extracts sample from RDD	
union(otherDataset):	unions two RDDs	
intersection(otherDataset):	finds intersection of two RDDs	
distinct([numPartitions]):	eliminates duplicates	
groupByKey([numPartitions]):	when called on RDD with (K, V) pairs, returns RDD with (K, Iterable) pairs	
sortByKey([ascending], [numPartitions]):	when called on (K, V) pairs where K implements Ordered ,	ISIT312 Big Data Management, SIM S4 2025

Spark Operations

Outline

The Programming Language Scala

Quick Start

Self Contained Application

Web User Interface

Operations on Resilient Distributed Datasets (RDDs)

Operations on Datasets

Operations on DataFrames

Spark SQL Module

Operations on Datasets

Operations on a **Dataset** start from creation of **case class**

```
case class Person(name: String, age: Long)  
defined class Person
```

Creating case class

```
val caseClassDS = Seq(Person("Andy", 32)).toDS()  
caseClassDS: org.apache.spark.sql.Dataset[Person] = [name: string, age: bigint]
```

Creating Dataset

```
caseClassDS.show()  
+---+---+  
| name | age |  
+---+---+  
| Andy | 32 |  
+---+---+
```

Listing Dataset

Dataset supports all operations of DataFrame

```
caseClassDS.select($"name").show()  
+---+  
| name |  
+---+  
| Andy |
```

Using a Dataset

Operations on Datasets

Operations on **Datasets** start from creation of **case class**

```
case class Flight(DEST_COUNTRY_NAME: String, ORIGIN_COUNTRY_NAME: String, count: BigInt)
```

Creating case class

Next we create a **DataFrame**

```
val flightsDF = spark.read.parquet("/mnt/defg/chapter-1-data/parquet/2010-summary.parquet")
```

Creating DataFrame

Finally, **DataFrame** is casted to **Dataset**

```
val flights = flightsDF.as[Flight]
```

Creating Dataset

Filtering a **Dataset**

```
def originIsDestination(flight_row: Flight): Boolean = {  
    return flight_row.ORIGIN_COUNTRY_NAME == flight_row.DEST_COUNTRY_NAME}
```

Defining a function

```
flights.filter(flight_row => originIsDestination(flight_row)).first()
```

Filtering

Mapping a **Dataset**

[TOP](#) val destinations = flights.map(f => f.DEST_COUNTRY_NAME)

Mapping
23/41

ISIT312 Big Data Management, SIM S4 2025

Spark Operations

Outline

The Programming Language Scala

Quick Start

Self Contained Application

Web User Interface

Operations on Resilient Distributed Datasets (RDDs)

Operations on Datasets

Operations on DataFrames

Spark SQL Module

Operations on DataFrames

Dataset and DataFrame are the data abstractions for Spark SQL

Dataset is a distributed collection of data

- It supports the use of self-defined functions to process data
- For example, `map` and `reduce` functions in the previous slides
- Dataset is typed; typing is checked at compiling time

DataFrame is a Dataset organized into named columns.

- It is conceptually equivalent to a table in a relational database or a data frame in R/Python
- To use self-defined functions, you need to register them with Spark
- DataFrame is untyped, i.e., typing is checked at runtime
- DataFrame is more performance-optimal than Dataset

Operations on DataFrames

DataFrame can be created in the following way

```
Creating a DataFrame  
val df = spark.read.json("people.json")  
df.show()
```

```
Results  
+---+-----+  
| age| name |  
+---+-----+  
| null| Michael|  
| 30 | Andy |  
| 19 | Justin|  
+---+-----+
```

```
df.printSchema()
```

```
Results  
root  
|-- age: long (nullable = true)  
|-- name: string (nullable = true)
```

Operations on DataFrames

Select on a DataFrame

```
df.select($"name", $"age" + 1).show()
```

Selecting from a DataFrame

```
+-----+-----+
| name | (age + 1) |
+-----+-----+
| Michael |      null |
| Andy    |        31 |
| Justin  |        20 |
+-----+-----+
```

Results

```
df.filter($"age" > 21).show()
```

Filtering a DataFrame

```
+---+---+
| age | name |
+---+---+
| 30 | Andy |
+---+---+
```

Results

Operations on DataFrames

Count people by age

```
df.groupBy("age").count().show()
```

Counting in a DataFrame

age	count
19	1
null	1
30	1

Results

Operations on DataFrames

Register a **DataFrame** as **SQL** temporary view

Registering and selecting from DataFrame

```
df.createOrReplaceTempView("people")
val sqlDF = spark.sql("SELECT * FROM people")
sqlDF.show()
```

Results

age	name
null	Michael
30	Andy
19	Justin

Operations on DataFrames

When to use **DataFrames** ?

Except for the following few cases, you can use them interchangeable (if performance is not a concern). You also can convert one to the other easily.

- In the Bigdata pipeline, you read an unstructured data source, for example, a text file as a **Dataset** and continue processing the data
- You can directly read an structured source like Hive table, JSON document as a **DataFrame**
- If you expect to use self-defined function easily, especially in the data cleaning or preprocessing stage of the pipeline, you should use a Dataset

Operations on DataFrames

Create a Dataset of `Person` objects from a text file and convert it to a DataFrame

Converting a Dataset to DataFrame

```
val peopleDF = spark.sparkContext  
.textFile("examples/src/main/resources/people.txt")  
.map(_.split(","))  
.map(attributes => Person(attributes(0), attributes(1).trim.toInt))  
.toDF()
```

Results

```
peopleDF: org.apache.spark.sql.DataFrame = [name: string, age: bigint]
```

Operations on DataFrames

Convert DataFrame to Dataset

Converting a Dataset to DataFrame

```
case class Employee(name: String, salary: Long)
val ds =
    spark.read.json(".../examples/src/main/resources/employees.json").as[Employee]
```

Results

```
ds: org.apache.spark.sql.Dataset[Employee] = [name: string, salary: bigint]
```

Operations on DataFrames

Spark DataFrame/Dataset support two types of operations:
transformations and actions

Transformations are operations on DataFrames/Datasets that return a new DataFrame/Dataset

- For example `select()`, `groupBy()`, `map()`, and `filter()`

Actions are operations that return a result to the driver program or write it to storage, and kick off a computation

- For example `show()`, `count()`, and `first()`

Return type difference: transformations return DataFrames/Datasets, whereas actions return some other data type

Spark treats the two operations very differently

Operations on DataFrames

Transformations are **lazily evaluated**, meaning that **Spark** will not begin to execute until it sees an action

Instead, **Spark** internally records metadata to indicate that some transformation operation has been requested

For example **transformation** creates another **DataFrame**

Creating a DataFrame

```
val sqlDF = spark.sql("SELECT * FROM people")
```

Action triggers the computation

Action on a DataFrame

```
sqlDF.show()
```

Operations on DataFrames

The **lazy evaluation** to reduce the number of passes it has to take over the dataset

In **Hadoop MapReduce**, developers often have to consider how to group together operations to minimize the number of MapReduce passes

In Spark, there is no substantial benefit to writing a single complex map instead of chaining together many simple operations

Thus, users are free to organize their program into smaller, more manageable operations

Spark Operations

Outline

The Programming Language Scala

Quick Start

Self Contained Application

Web User Interface

Operations on Resilient Distributed Datasets (RDDs)

Operations on Datasets

Operations on DataFrames

Spark SQL

Spark SQL

Spark SQL is a Spark module for general data processing and analytics. It can be used for all sorts of data, from unstructured log files to semi-structured CSV files and highly structured Parquet files.

To interact with Spark SQL, you can either use SQL or Spark Structured API, or both.

The same execution engine is used, independent of which API/language you use to express the computation.

The APIs of Spark SQL provide a rich set of pre-built, high-level operations for accomplishing sophisticated data processing and ETL jobs, and mechanism to implement your own operations, for example self-defined functions and aggregations.

Spark SQL

Spark SQL has two data abstractions

- DataFrame
- Dataset (available in Scala/Java APIs, but not Python/R APIs)
- DataFrame can be represented as SQL tables and views

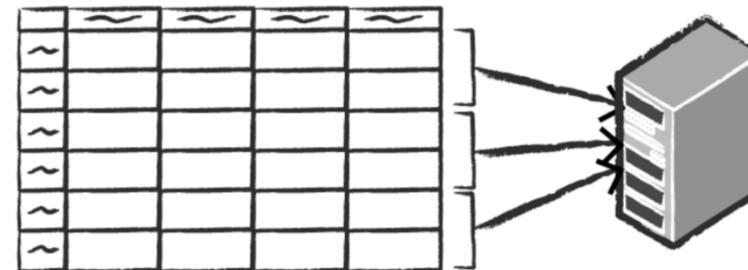
Both are distributed table-like collections with well-defined rows and columns.

- DataFrame vs. spreadsheet

Spreadsheet on
a single machine



Table or Data Frame
partitioned across servers
in a data center



Spark SQL

Spark SQL allows to code **SQL** statements in **Scala**, **Java** and **Python** language APIs.

To use **SQL** to manipulate a **DataFrame**, we first need to create a temporal view for it

Creating a temporal view

```
df.createOrReplaceTempView("dfTable")
```

All standard **SQL** statements + functions are applicable in **Spark SQL**

Spark implements a subset of [ANSI SQL:2003](#)

Spark SQL

Using SQL

Applying sql method

```
spark.sql(  
    "SELECT DEST_COUNTRY_NAME, sum(count)  
     FROM dfTable  
    GROUP BY DEST_COUNTRY_NAME"  
)  
.where("DEST_COUNTRY_NAME like 'S%'")  
.where("sum(count) > 10")  
.show(2)
```

Results

DEST_COUNTRY_NAME	sum(count)
Senegal	40
Sweden	118

References

The Scala Programming Language

A Gentle Introduction to Spark, Databricks, (Available in **READINGS** folder)

RDD Programming Guide

Spark SQL, DataFrames and Datasets Guide

Karau H., Fast data processing with Spark Packt Publishing, 2013
(Available from UOW Library)

Srinivasa, K.G., Guide to High Performance Distributed Computing: Case Studies with Hadoop, Scalding and SparkSpringer, 2015 (Available from UOW Library)

Chambers B., Zaharia M., Spark: The Definitive Guide, O'Reilly 2017

Perrin J-G., Spark in Action, 2nd ed., Manning Publications Co. 2020